

Practical 3 - BGN 54

October 28, 2023

```
[3]: from scipy.stats import f_oneway, normaltest, levene
import pandas as pd
import seaborn as sns
import numpy as np
from IPython.display import display, HTML
import statsmodels.formula.api as smf

sns.set_theme(style="ticks", palette="pastel")
```

1 Part I - Rubisco transgene experiment

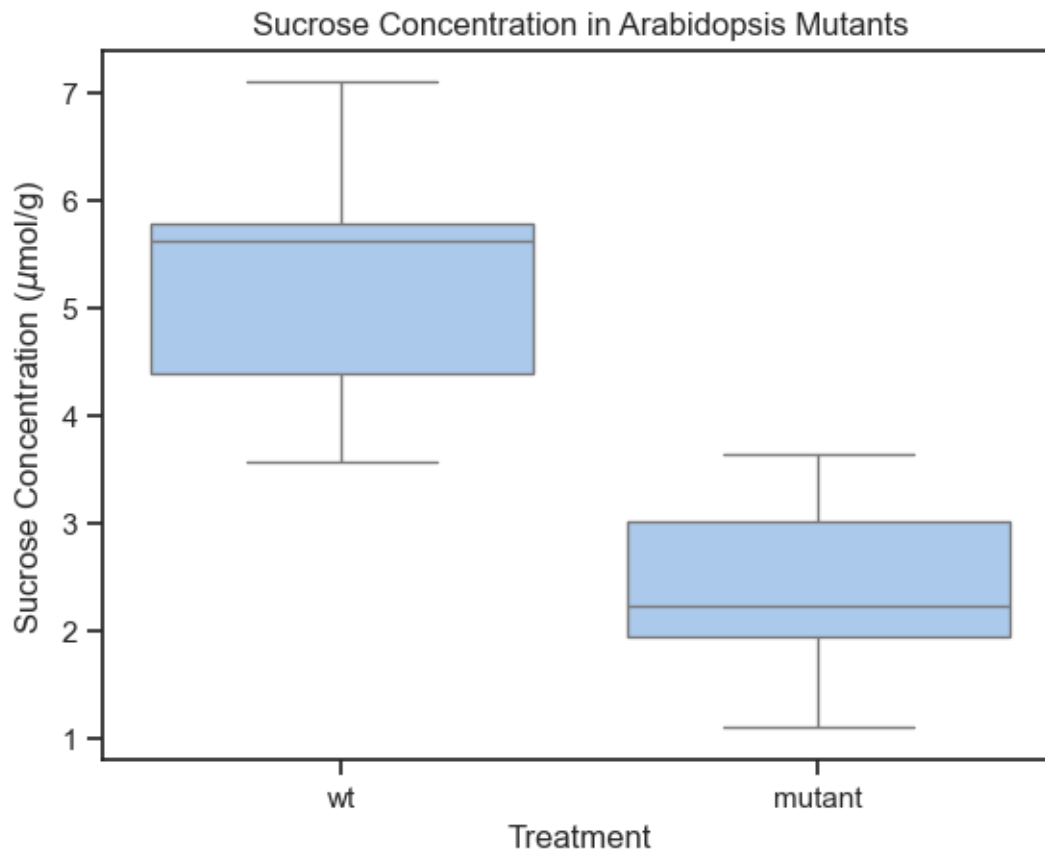
1.1 Sucrose

```
[4]: sucrose_data = pd.read_csv("data/Sucrose.csv") # read the data from the csv file
display(sucrose_data)
# create the boxplot
ax = sns.boxplot(x="treatment", y="Sucrose", data=sucrose_data)
ax.set_title("Sucrose Concentration in Arabidopsis Mutants")
ax.set_xlabel("Treatment")
ax.set_ylabel("Sucrose Concentration ($\mu$mol/g)")
```

	Sucrose	treatment
0	5.68	wt
1	7.10	wt
2	5.01	wt
3	3.57	wt
4	3.78	wt
5	5.63	wt
6	6.63	wt
7	3.76	wt
8	5.65	wt
9	5.63	wt
10	5.90	wt
11	1.10	mutant
12	2.07	mutant
13	1.84	mutant
14	3.03	mutant
15	3.48	mutant

16	1.55	mutant
17	3.65	mutant
18	2.81	mutant
19	2.24	mutant
20	2.09	mutant
21	3.00	mutant

```
[4]: Text(0, 0.5, 'Sucrose Concentration ( $\mu\text{mol/g}$ )')
```



In order to do a one-way ANOVA, we should test that the data collected fits the assumptions of the test. The assumptions of a one-way ANOVA are:

- The samples are independent.
- Each sample is from a normally distributed population. This can be tested using a SciPy's `normaltest` (It is based on D'Agostino and Pearson's test that combines skew and kurtosis to produce an omnibus test of normality <https://docs.scipy.org/doc/scipy/reference/generated/scipy.stats.normaltest.html>).
- The SD of each group is the same (homogeneity of variance). This can be tested using the Levene's test.

Using the `normaltest` function:

```
[5]: # test if the distributions of each population is normally distributed

print(normaltest(sucrose_data[sucrose_data["treatment"] == "wt"]["Sucrose"]))
print(normaltest(sucrose_data[sucrose_data["treatment"] == "mutant"]
    ["Sucrose"]))
```

```
NormaltestResult(statistic=0.5832213517389794, pvalue=0.7470593275867403)
NormaltestResult(statistic=0.6009762314553877, pvalue=0.7404567038950335)

/Users/th/.pyenv/versions/3.12.0/lib/python3.12/site-packages/scipy/stats/_stats_py.py:1806: UserWarning: kurtosistest only valid for
n>=20 ... continuing anyway, n=11
  warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
```

This returns a high p-value, meaning that the probability of this data being normally distributed is high.

Using the Levene's test to test for homogeneity of variance:

```
[6]: print(levene(sucrose_data[sucrose_data["treatment"] == "wt"]["Sucrose"],
    sucrose_data[sucrose_data["treatment"] == "mutant"]["Sucrose"]))
```

```
LeveneResult(statistic=0.35140479282257187, pvalue=0.5599628035141764)
```

The high p-value means that the probability of this data having the same variance is high.

With these two assumptions verified we are able to do a one-way ANOVA.

```
[7]: (sucrose_f := f_oneway(sucrose_data[sucrose_data["treatment"] == "wt"]
    ["Sucrose"], sucrose_data[sucrose_data["treatment"] == "mutant"]
    ["Sucrose"]))
```

```
[7]: F_onewayResult(statistic=44.590599436650116, pvalue=1.6849610967522526e-06)
```

We were able to conclude that there was a difference in sucrose concentration between the wild type treatment and mutant treatment based on a one-way ANOVA ($F_{df=2,19} = 44.59, p < 0.0001$)

1.2 ACi

```
[8]: aci_df = pd.read_csv("data/stackedACi.csv")
display(aci_df)
```

	Treatment	Ci	A	x
0	wt	73.296105	-0.012118	0.067692
1	wt	116.519247	1.767882	0.137241
2	wt	147.280515	2.779511	0.180736
3	wt	184.439371	4.480535	0.227764
4	wt	235.531784	5.727014	0.284256
5	wt	286.786454	7.199071	0.333190
6	mutant	76.139613	-0.976707	0.072611
7	mutant	154.959330	-0.279311	0.190918

```

8      mutant  201.514984 -0.179819  0.247611
9      mutant  255.052613  0.280199  0.303716
10     mutant  323.141938  0.586070  0.364031
11     mutant  386.616168  0.997026  0.411550

```

```

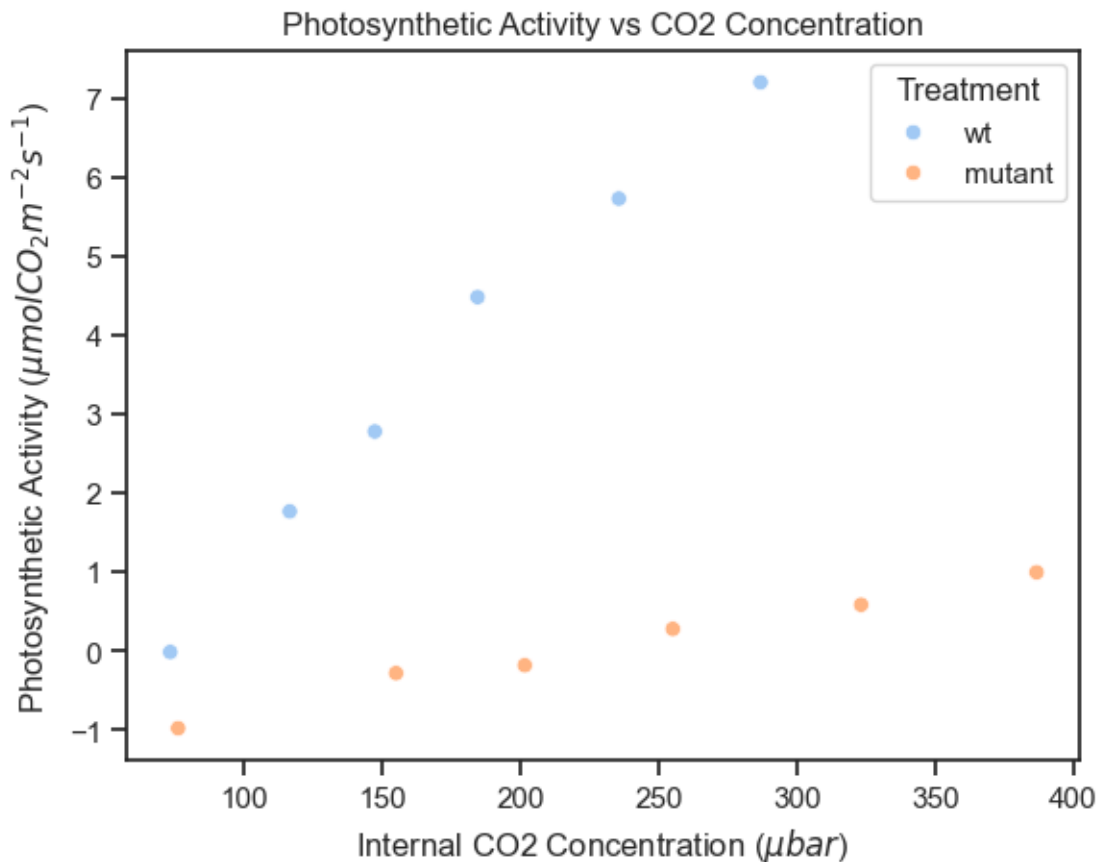
[9]: ax = sns.scatterplot(data=aci_df, x="Ci", y="A", hue="Treatment")
      ax.set_title("Photosynthetic Activity vs CO2 Concentration")
      ax.set_xlabel("Internal CO2 Concentration ( $\mu\text{bar}$ )")
      ax.set_ylabel("Photosynthetic Activity ( $\mu\text{mol CO}_2\text{ m}^{-2}\text{s}^{-1}$ )")

```

```

[9]: Text(0, 0.5, 'Photosynthetic Activity ( $\mu\text{mol CO}_2\text{ m}^{-2}\text{s}^{-1}$ )')

```



```

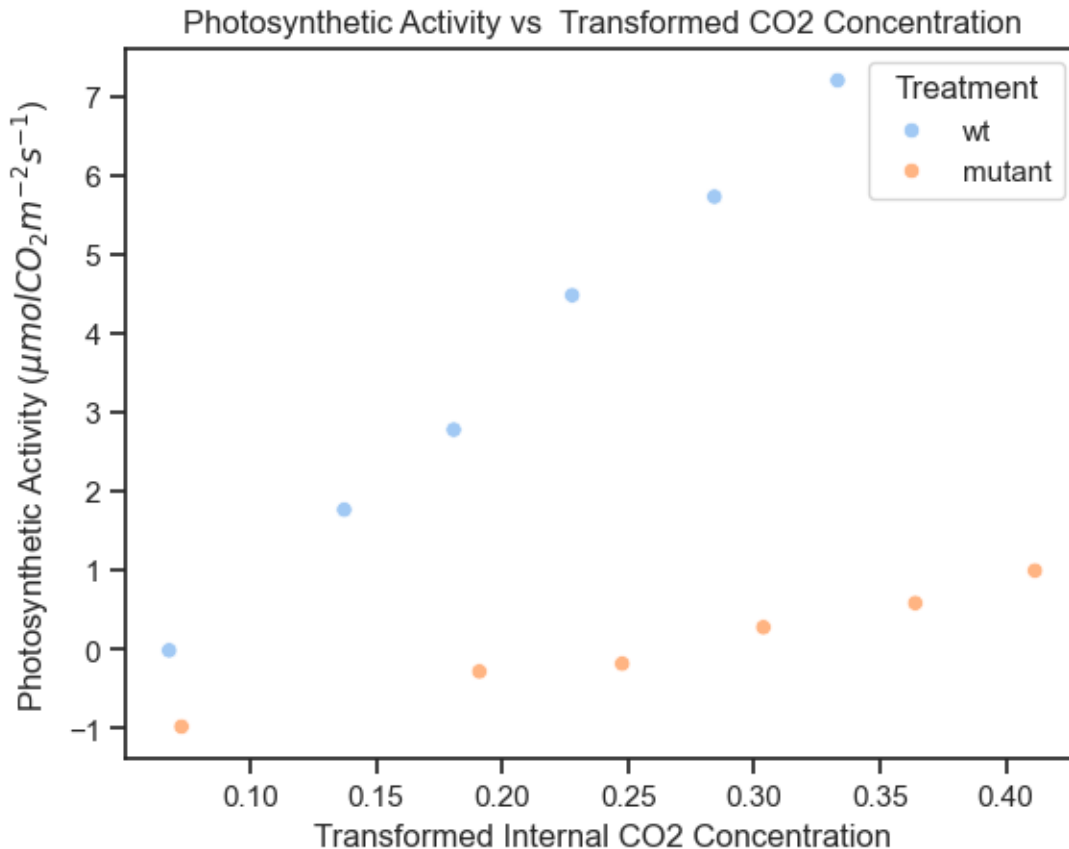
[10]: ax = sns.scatterplot(data=aci_df, x="x", y="A", hue="Treatment")
       ax.set_title("Photosynthetic Activity vs Transformed CO2 Concentration")
       ax.set_xlabel("Transformed Internal CO2 Concentration")
       ax.set_ylabel("Photosynthetic Activity ( $\mu\text{mol CO}_2\text{ m}^{-2}\text{s}^{-1}$ )")

```

```

[10]: Text(0, 0.5, 'Photosynthetic Activity ( $\mu\text{mol CO}_2\text{ m}^{-2}\text{s}^{-1}$ )')

```



```
[11]: # run levene's test on the transformed data
print(levene(aci_df[aci_df["Treatment"] == "wt"]["x"],
             aci_df[aci_df["Treatment"] == "mutant"]["x"])))
```

```
LeveneResult(statistic=0.28469435872023646, pvalue=0.6052965796161142)
```

The above shows that the levene test is not significant, meaning that the variance is the same between the two groups.

Both datasets appear linear, however will now generate a model and test whether and how the response of photosynthetic activity to CO2 concentration is affected by the transgene.

```
[12]: # create the predictors and response variables
aci_predictor = aci_df[["x", "Treatment"]]
aci_response = aci_df[["A"]]
# aci_predictor = pd.get_dummies(aci_predictor, drop_first=True) # create the
# dummy variables

# aci_predictor = smf.add_constant(aci_predictor) # add the constant term

display(aci_predictor)
```

		x Treatment
0	0.067692	wt
1	0.137241	wt
2	0.180736	wt
3	0.227764	wt
4	0.284256	wt
5	0.333190	wt
6	0.072611	mutant
7	0.190918	mutant
8	0.247611	mutant
9	0.303716	mutant
10	0.364031	mutant
11	0.411550	mutant

```
[13]: # fit the linear model
aci_linear_model = smf.ols(formula="A ~ x*C(Treatment)", data=aci_df).fit()
```

```
[14]: display(aci_linear_model.summary())

# print the p values for the interaction terms
print(aci_linear_model.pvalues["x:C(Treatment)[T.wt]"])

ax = sns.lmplot(data=aci_df, x="x", y="A", hue="Treatment")
ax.set_axis_labels("Transformed Internal CO2 Concentration", "Photosynthetic_
↳Activity ($\mu$ mol CO_2 m^{-2}s^{-1}$)")
```

```
/Users/th/.pyenv/versions/3.12.0/lib/python3.12/site-
packages/scipy/stats/_stats_py.py:1806: UserWarning: kurtosistest only valid for
n>=20 ... continuing anyway, n=12
warnings.warn("kurtosistest only valid for n>=20 ... continuing ")
```

Dep. Variable:	A	R-squared:	0.998
Model:	OLS	Adj. R-squared:	0.997
Method:	Least Squares	F-statistic:	1386.
Date:	Sat, 28 Oct 2023	Prob (F-statistic):	3.34e-11
Time:	16:27:42	Log-Likelihood:	9.4037
No. Observations:	12	AIC:	-10.81
Df Residuals:	8	BIC:	-8.868
Df Model:	3		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-1.4251	0.142	-10.056	0.000	-1.752	-1.098
C(Treatment)[T.wt]	-0.5167	0.199	-2.601	0.032	-0.975	-0.059
x	5.6451	0.492	11.466	0.000	4.510	6.780
x:C(Treatment)[T.wt]	21.6466	0.794	27.272	0.000	19.816	23.477


Omnibus:	0.184	Durbin-Watson:	2.842
Prob(Omnibus):	0.912	Jarque-Bera (JB):	0.167
Skew:	-0.187	Prob(JB):	0.920
Kurtosis:	2.559	Cond. No.	26.7

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

3.521734306130354e-09

[14]: <seaborn.axisgrid.FacetGrid at 0x294be0c80>



From this data table, the p-value for the interaction term is $p > 0.001$ therefore we cannot drop the interaction term.

From the table we can see:

- The intercept (mutant without wt) is -1.42,
- The slope is normally 5.645 for the mutant
- The wt leads to a decrease of intercept by -0.51
- The wt leads to an increase in slope of 21.64

This means the equation for the mutant is: $A = -1.42 + 5.645x$ and the equation for the wt is: $y = -1.93 + 27.285x$.

From the above, we can see that the Rd (the negative intercept), the positive is the the rate of CO₂ release from mitochondrial respiration, is $1.42 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$ in the mutant line and $1.93 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$ in the wild type.

We can see that the maximum rate of RuBP carboxylation (V_{max}, which is the gradient) is dependent on the treatment. For the mutant, V_{max} is $21.6 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$, and for the wild type, V_{max} is $27.2 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$.

Summarising in a statement that could be found in a paper:

The mutant line had a lower rate of CO₂ release from mitochondrial respiration ($1.42 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$) than the wild type ($1.93 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$). The maximum rate of RuBP carboxylation (V_{max}) was higher in the wild type ($27.2 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$) than the mutant ($21.6 \mu \text{ mol CO}_2 \text{ m}^{-2}\text{s}^{-1}$) (OLS regression ANOVA, $F_{df=3,8} = 1386, p < 0.0001$).

2 Part II - Past Part IB examination question

- a) What is being done in the analyses?

You are testing to see if it is possible to generate a linear model that generates the predicted yield based on the treatment (virus) type the plant is exposed to. You are then showing a summary, which shows the intercept (virus cc), and the effect that the three other viruses have relative to this; it is also showing the p-value for the model, which is significant, meaning that the model is a good fit for the data. This is followed by the ANOVA which shows the F test completed to gain the p-value stated in the summary. The last line shows the effect on the p-value. There is also a difference between the two analyses is that the first analysis set is done on the raw data, whereas the second analysis is done on the log transformed data.

- b) Which of the two analyses do you think is most appropriate, and why?

From the diagnostic plots, although the normal Q-Q plot has a small tail on the end off the 1:1 (point 1), the fit to the normal Q-Q graph is much better on the log modified data. Also, the data is much more evenly spread on the residual vs fitted plot. Therefore, I think that the log modified data is more appropriate. This is confirmed by the lower p-value for the log modified data.

- c) What can be concluded from your preferred analysis?

Compared to the control (oo) treatment, all three treatments affected the plant, however the feathery mottle virus treatment (ff) affected the log(yield) of the plant the least, followed by the chlorotic dwarf treatment (cc), with the most effect coming from the combined treatments (fc). This is tested with the ANOVA test which resulted in a p-value of approximately 0.00015, which is significant, meaning that the model is a good fit for the data ($F_{df=3,8} = 26.4$).

- d) Write a sentence or two – in the style of the results section of a scientific publication – showing how the analysis could be reported.

The yield of sweetpotato was significantly affected by the virus treatments ($F_{df=3,8} = 26.4, p < 0.0005$). The feathery mottle virus treatment (ff) affected the log(yield) of the plant the least, followed by the chlorotic dwarf treatment (cc), with the most effect coming from the combined treatments (fc).

e) Estimate the total yield of sweetpotato in each of the three plots containing:

part i) sweetpotato with chlorotic dwarf virus (i.e. the treatment cc);

From the model, the fitted value was 1.38..., and using residuals vs fitted plot, at around a fitted value of 1.38, the three plots had residuals of 0.075, -0.025 and -0.05. Therefore the three plots yield (kg) can be calculated:

```
[15]: np.power(10, 1.38+np.array([0.075, -0.025, -0.05]))
```

```
[15]: array([28.51018268, 22.64644308, 21.3796209 ])
```

part ii) doubly-infected sweetpotato plants (i.e. the treatment fc).

Using a similar method, with the fitted value for fc being $1.38 - 0.27 = 1.11$, the three plots had residuals of 0.07, 0.01 and -0.08. Therefore the three plots yield (kg) can be calculated:

```
[16]: np.power(10, 1.38-0.27+np.array([0.07, -0.01, -0.08]))
```

```
[16]: array([15.13561248, 12.58925412, 10.71519305])
```